

Task-D: Collinear features and their effect on linear models

```
In [55]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import GridSearchCV
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn.linear_model
```

```
In [56]: data = pd.read_csv('task_d.csv')
```

```
In [57]: data.head()
```

```
Out[57]:
```

	x	y	z	x*x	2*y	2*z+3*x*x	w	target
0	-0.581066	0.841837	-1.012978	-0.604025	0.841837	-0.665927	-0.536277	0
1	-0.894309	-0.207835	-1.012978	-0.883052	-0.207835	-0.917054	-0.522364	0
2	-1.207552	0.212034	-1.082312	-1.150918	0.212034	-1.166507	0.205738	0
3	-1.364174	0.002099	-0.943643	-1.280666	0.002099	-1.266540	-0.665720	0
4	-0.737687	1.051772	-1.012978	-0.744934	1.051772	-0.792746	-0.735054	0

```
In [58]: X = data.drop(['target'], axis=1).values
Y = data['target'].values
#print(X)
```

Doing perturbation test to check the presence of collinearity

Task: 1 Logistic Regression

- Finding the Correlation between the features**
 - check the correlation between the features
 - plot heat map of correlation matrix using seaborn heatmap
- Finding the best model for the given data**
 - Train Logistic regression on data(X,Y) that we have created in the above cell
 - Find the best hyper parameter alpha with hyper parameter tuning using k-fold cross validation (grid search CV or random search CV make sure you choose the alpha in log space)
 - Create a new Logistic regression with the best alpha (search for how to get the best hyper parameter value), name the best model as 'best_model'
- Getting the weights with the original data**
 - train the 'best_model' with X, Y
 - Check the accuracy of the model 'best_model_accuracy'
 - Get the weights W using best_model.coef_
- Modifying original data**
 - Add a noise (order of 10^{-2}) to each element of X and get the new data set X' ($X' = X + e$)
 - Train the same 'best_model' with data (X', Y)
 - Check the accuracy of the model 'best_model_accuracy_edited'
 - Get the weights W' using best_model.coef_
- Checking deviations in metric and weights**
 - find the difference between 'best_model_accuracy_edited' and 'best_model_accuracy'
 - find the absolute change between each value of W and W' ==> $|W - W'|$
 - print the top 4 features which have higher % change in weights compare to the other feature

Task: 2 Linear SVM

- Do the same steps (2, 3, 4, 5) we have done in the above task 1.

Do write the observations based on the results you get from the deviations of weights in both Logistic Regression and linear SVM

Task: 1 Logistic Regression

```
In [59]: data.drop('target', axis = 1, inplace = True)
```

```
In [60]: # 1. Finding the Correlation between the features  
correlations = data.corr()
```


```
In [61]: correlations
```

```
Out[61]:
```

	x	y	z	x*x	2*y	2*z+3*x*x	w
x	1.000000	-0.205926	0.812458	0.997947	-0.205926	0.996252	0.583277
y	-0.205926	1.000000	-0.602663	-0.209289	1.000000	-0.261123	-0.401790
z	0.812458	-0.602663	1.000000	0.807137	-0.602663	0.847163	0.674486
x*x	0.997947	-0.209289	0.807137	1.000000	-0.209289	0.997457	0.583803
2*y	-0.205926	1.000000	-0.602663	-0.209289	1.000000	-0.261123	-0.401790
2*z+3*x*x	0.996252	-0.261123	0.847163	0.997457	-0.261123	1.000000	0.606860
w	0.583277	-0.401790	0.674486	0.583803	-0.401790	0.606860	1.000000

```
In [62]: sns.heatmap(correlations)
```

```
Out[62]: <AxesSubplot:>
```



```
In [63]: clf = sklearn.linear_model.SGDClassifier(max_iter=1000, loss = 'log')  
Cs = np.logspace(-5,5,11)  
tuned_parameters = [{'alpha': Cs}]  
model = GridSearchCV(clf, tuned_parameters, scoring = 'accuracy', cv=4)  
model.fit(X, Y)
```

```
Out[63]: GridSearchCV(cv=4, estimator=SGDClassifier(loss='log'),  
param_grid=[{'alpha': array([1.e-05, 1.e-04, 1.e-03, 1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+02,  
1.e+03, 1.e+04, 1.e+05])}],  
scoring='accuracy')
```

```
In [64]: #2.finding best hyper parameter  
best_param = model.best_params_  
print(best_param)  
  
{'alpha': 0.0001}
```

```
In [65]: # 3. Getting the weights with the original data  
best_model = sklearn.linear_model.SGDClassifier(alpha = 1e-05, loss = 'log')  
best_model.fit(X,Y)  
best_model_accuracy = best_model.score(X, Y)  
print(best_model_accuracy)  
W = best_model.coef_  
print(W)  
  
1.0  
[[ 12.03098932 -23.27128945  23.9332139   14.31455572 -23.27128945  
  15.77274148  13.19109547]]
```

```
In [67]: #4. Modifying original data  
# a. Add a noise(order of 10^-2) to each element of X and get the new data set X' (X' = X + e)  
df = pd.read_csv('task_d.csv')
```

```
df.drop(['target'], axis = 1, inplace = True)
df.applymap(lambda x: x + 0.01)
X_noise = df.values
```

```
In [69]: # b. Train the same 'best_model' with data (X', Y)
best_model = sklearn.linear_model.SGDClassifier(alpha = 1e-05, loss = 'log')
best_model.fit(X_noise, Y)
best_model_accuracy_edited = best_model.score(X_noise, Y)
print(best_model_accuracy_edited)
W_noise = best_model.coef_
print(W_noise)
```

```
1.0
[[ 21.71700043 -28.9284769  37.41549186  20.74322511 -28.9284769
  23.18634811   8.14267415]]
```

```
In [70]: #5. Checking deviations in metric and weights
```

```
In [71]: #a. find the difference between 'best_model_accuracy_edited' and 'best_model_accuracy'
best_model_accuracy_edited - best_model_accuracy
```

```
Out[71]: 0.0
```

```
In [72]: #b. find the absolute change between each value of W and W' ==> |(W-W')|
abs = np.abs(W - W_noise)
abs
```

```
Out[72]: array([[ 9.68601111,  5.65718746, 13.48227796,  6.42866939,  5.65718746,
  7.41360662,  5.04842132]])
```

```
In [73]: l = sorted(abs[0], reverse = True)
cols = data.columns
cols
```

```
Out[73]: Index(['x', 'y', 'z', 'x*x', '2*y', '2*z+3*x*x', 'w'], dtype='object')
```

```
In [74]: # c. print the top 4 features which have higher % change in weights compare to the other feature
def Top_4_feature(abs):
    list = []
    for i in range(4):
        idx = abs.argmax()
        print(idx)
        list.append('Feature Name ' + cols[idx] + ' : ' + str(l[i]))
        abs = np.delete(abs, idx)
    return list
```

```
In [75]: Top_4_feature(abs)
```

```
2
0
3
1
```

```
Out[75]: ['Feature Name z : 13.482277956714597',
'Feature Name x : 9.686011113345108',
'Feature Name x*x : 7.4136066229736155',
'Feature Name y : 6.428669387808686']
```

Task: 2 Linear SVM

```
In [40]: clf = sklearn.linear_model.SGDClassifier(max_iter=1000, loss = 'hinge')
Cs = np.logspace(-5,5,11)
tuned_parameters = [{'alpha': Cs}]
model = GridSearchCV(clf, tuned_parameters, scoring = 'accuracy', cv=4)
model.fit(X, Y)
```

```
Out[40]: GridSearchCV(cv=4, estimator=SGDClassifier(),
    param_grid=[{'alpha': array([1.e-05, 1.e-04, 1.e-03, 1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+02,
    1.e+03, 1.e+04, 1.e+05])}],
    scoring='accuracy')
```

```
In [41]: #2.finding best hyper parameter
best_param = model.best_params_
print(best_param)

{'alpha': 1e-05}
```

```
In [42]: # 3. Getting the weights with the original data
best_model = sklearn.linear_model.SGDClassifier(alpha = 1e-05, loss = 'hinge')
best_model.fit(X,Y)
best_model_accuracy = best_model.score(X, Y)
print(best_model_accuracy)
W = best_model.coef_
print(W)

1.0
[[ 17.77791249 -19.57998183  37.12646288  14.98991774 -19.57998183
   17.97342313 -14.69082429]]
```

```
In [43]: # b. Train the same 'best_model' with data (X', Y)
best_model = sklearn.linear_model.SGDClassifier(alpha = 1e-05, loss = 'hinge')
best_model.fit(X_noise,Y)
best_model_accuracy_edited = best_model.score(X_noise, Y)
print(best_model_accuracy_edited)
W_noise = best_model.coef_
print(W_noise)

1.0
[[ 27.80684591 -28.9035864   64.61034056  22.61125655 -28.9035864
   28.15088679  24.33218649]]
```

```
In [44]: #5. Checking deviations in metric and weights
```

```
In [45]: #a. find the difference between 'best_model_accuracy_edited' and 'best_model_accuracy'
best_model_accuracy_edited - best_model_accuracy
```

```
Out[45]: 0.0
```

```
In [46]: #b. find the absolute change between each value of W and W' ==> |(W-W')|
abs = np.abs(W - W_noise)
abs
```

```
Out[46]: array([[10.02893343,  9.32360457, 27.48387767,  7.6213388 ,  9.32360457,
   10.17746365, 39.02301078]])
```

```
In [47]: l = sorted(abs[0], reverse = True)
cols = data.columns
cols
```

```
Out[47]: Index(['x', 'y', 'z', 'x*x', '2*y', '2*z+3*x*x', 'w'], dtype='object')
```

```
In [48]: # c. print the top 4 features which have higher % change in weights compare to the other feature
def Top_4_feature(abs):
    list = []
    for i in range(4):
        idx = abs.argmax()
        list.append('Feature Name ' + cols[idx] + ' : ' + str(l[i]))
        abs = np.delete(abs, idx)
    return list
```

```
In [49]: Top_4_feature(abs)
```

```
Out[49]: ['Feature Name w : 39.02301078138523',
 'Feature Name z : 27.483877674389305',
 'Feature Name 2*y : 10.177463653563741',
 'Feature Name x : 10.02893342611118']
```

Observation Task 1: deviations on weights(Logistic regression) a) more deviations when colinearity between the features are high otherwise only a min deviation b) executing the model again and again based on the weight deviation and colinearity the values should be increased c) Implementation done on

Logistic loss Task 2: deviations on weights(SVM) a) more deviations when colinearity between the features are high otherwise only a min deviation b) executing the model again and again based on the weight deviation and colinearity the values should be increased c) Implementation done on hinge loss

In []:

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js